

transposition

April 1, 2025

```
[1]: import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import librosa
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import utils
from sklearn.calibration import LabelEncoder
```

```
[2]: tracks = utils.load('data/fma_metadata/tracks.csv')
genres = utils.load('data/fma_metadata/genres.csv')
features = utils.load('data/fma_metadata/features.csv')
small = tracks[tracks['set', 'subset'] <= 'small']
```

```
[3]: print(tracks.head())
```

	album							
	comments		date_created	date_released	engineer	favorites	id	
track_id								
2	0	2008-11-26 01:44:45	2009-01-05	NaN	4	1		
3	0	2008-11-26 01:44:45	2009-01-05	NaN	4	1		
5	0	2008-11-26 01:44:45	2009-01-05	NaN	4	1		
10	0	2008-11-26 01:45:08	2008-02-06	NaN	4	6		
20	0	2008-11-26 01:45:05	2009-01-06	NaN	2	4		

	information	listens	producer	tags	
track_id					
2	<p></p>	6073	NaN	[]	
3	<p></p>	6073	NaN	[]	
5	<p></p>	6073	NaN	[]	
10	NaN	47632	NaN	[]	
20	<p> "spiritual songs" from Nicky Cook</p>	2710	NaN	[]	

	...	track		\
	...	information	interest	language_code
track_id	...			
2	...	NaN	4656	en
3	...	NaN	1470	en
5	...	NaN	1933	en
10	...	NaN	54881	en
20	...	NaN	978	en

		license	listens	lyricist
track_id				
2	Attribution-NonCommercial-ShareAlike 3.0 Inter...	1293	NaN	
3	Attribution-NonCommercial-ShareAlike 3.0 Inter...	514	NaN	
5	Attribution-NonCommercial-ShareAlike 3.0 Inter...	1151	NaN	
10	Attribution-NonCommercial-NoDerivatives (aka M...	50135	NaN	
20	Attribution-NonCommercial-NoDerivatives (aka M...	361	NaN	

	number	publisher	tags	title
track_id				
2	3	NaN	[]	Food
3	4	NaN	[]	Electric Ave
5	6	NaN	[]	This World
10	1	NaN	[]	Freeway
20	3	NaN	[]	Spiritual Level

[5 rows x 52 columns]

```
[4]: # Load the features and metadata
features_df = pd.read_csv('data/fma_metadata/features.csv', index_col=0)
tracks_df = pd.read_csv('data/fma_metadata/tracks.csv', header=[0, 1],
    ↪ index_col=0)

# Flatten the multi-level column names in tracks.csv
tracks_df.columns = ['_'.join(col).strip() if isinstance(col, tuple) else col,
    ↪ for col in tracks_df.columns]

# Filter the small dataset (subset)
small_tracks = tracks_df[tracks_df['set_subset'] == 'small']

# Extract relevant columns: track_id, features, and genre
small_tracks = small_tracks[['track_genre_top']].copy()
small_tracks.index = small_tracks.index.astype(int) # Ensure track_id is an
    ↪ integer
```

```
# Merge features with genres
mapping_df = features_df.merge(small_tracks, left_index=True, right_index=True)

# Save the mapping to a new CSV file
mapping_df.to_csv('features_genre_mapping_small.csv', index=True)

print("Mapping of features to genres for the small dataset saved to_
↳ 'features_genre_mapping_small.csv'.")
```

```
C:\Users\james\AppData\Local\Temp\ipykernel_17884\3966755446.py:2: DtypeWarning:
Columns (0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19,20,21,22,23,24,25,26,
27,28,29,30,31,32,33,34,35,36,37,38,39,40,41,42,43,44,45,46,47,48,49,50,51,52,53
,54,55,56,57,58,59,60,61,62,63,64,65,66,67,68,69,70,71,72,73,74,75,76,77,78,79,8
0,81,82,83,84,85,86,87,88,89,90,91,92,93,94,95,96,97,98,99,100,101,102,103,104,1
05,106,107,108,109,110,111,112,113,114,115,116,117,118,119,120,121,122,123,124,1
25,126,127,128,129,130,131,132,133,134,135,136,137,138,139,140,141,142,143,144,1
45,146,147,148,149,150,151,152,153,154,155,156,157,158,159,160,161,162,163,164,1
65,166,167,168,169,170,171,172,173,174,175,176,177,178,179,180,181,182,183,184,1
85,186,187,188,189,190,191,192,193,194,195,196,197,198,199,200,201,202,203,204,2
05,206,207,208,209,210,211,212,213,214,215,216,217,218,219,220,221,222,223,224,2
25,226,227,228,229,230,231,232,233,234,235,236,237,238,239,240,241,242,243,244,2
45,246,247,248,249,250,251,252,253,254,255,256,257,258,259,260,261,262,263,264,2
65,266,267,268,269,270,271,272,273,274,275,276,277,278,279,280,281,282,283,284,2
85,286,287,288,289,290,291,292,293,294,295,296,297,298,299,300,301,302,303,304,3
05,306,307,308,309,310,311,312,313,314,315,316,317,318,319,320,321,322,323,324,3
25,326,327,328,329,330,331,332,333,334,335,336,337,338,339,340,341,342,343,344,3
45,346,347,348,349,350,351,352,353,354,355,356,357,358,359,360,361,362,363,364,3
65,366,367,368,369,370,371,372,373,374,375,376,377,378,379,380,381,382,383,384,3
85,386,387,388,389,390,391,392,393,394,395,396,397,398,399,400,401,402,403,404,4
05,406,407,408,409,410,411,412,413,414,415,416,417,418,419,420,421,422,423,424,4
25,426,427,428,429,430,431,432,433,434,435,436,437,438,439,440,441,442,443,444,4
45,446,447,448,449,450,451,452,453,454,455,456,457,458,459,460,461,462,463,464,4
65,466,467,468,469,470,471,472,473,474,475,476,477,478,479,480,481,482,483,484,4
85,486,487,488,489,490,491,492,493,494,495,496,497,498,499,500,501,502,503,504,5
05,506,507,508,509,510,511,512,513,514,515,516,517,518) have mixed types.
```

Specify dtype option on import or set low_memory=False.

```
features_df = pd.read_csv('data/fma_metadata/features.csv', index_col=0)
```

Mapping of features to genres for the small dataset saved to
'features_genre_mapping_small.csv'.

```
[5]: mapping_df.describe()
```

```
[5]:
```

	chroma_cens	chroma_cens.1	chroma_cens.2	chroma_cens.3	\
count	7918.000000	7918.0000	7918.000000	7918.000000	
unique	7901.000000	7901.0000	7901.000000	7900.000000	
top	2.120035	-0.1481	0.687363	1.388522	
freq	6.000000	6.0000	6.000000	6.000000	

	chroma_cens.4	chroma_cens.5	chroma_cens.6	chroma_cens.7	\
count	7918.000000	7918.000000	7918.000000	7918.000000	
unique	7900.000000	7900.000000	7900.000000	7901.000000	
top	-0.779828	-1.049356	-1.352663	-1.146229	
freq	6.000000	6.000000	6.000000	6.000000	

	chroma_cens.8	chroma_cens.9	...	tonnetz.40	tonnetz.41	\
count	7918.000000	7918.000000	...	7918.000000	7918.000000	
unique	7901.000000	7901.000000	...	7899.000000	7900.000000	
top	0.13384	0.494833	...	0.023668	0.022414	
freq	6.000000	6.000000	...	6.000000	6.000000	

	zcr	zcr.1	zcr.2	zcr.3	zcr.4	\
count	7918.000000	7918.000000	7918.000000	7918.000000	7918.0	
unique	7901.000000	1573.000000	7900.000000	343.000000	85.0	
top	101.444832	0.385742	0.040447	0.033691	0.0	
freq	6.000000	20.000000	6.000000	99.000000	3648.0	

	zcr.5	zcr.6	track_genre_top
count	7918.000000	7918.000000	7918
unique	7900.000000	7897.000000	8
top	5.669625	0.026036	Electronic
freq	6.000000	6.000000	1000

[4 rows x 519 columns]

```
[6]: # Filter relevant features: mfcc_* and spectral_*
# feature_columns = [col for col in mapping_df.columns if col.startswith('mfcc.'
# ↪) or col.startswith('spectral_')]
feature_columns = [col for col in mapping_df.columns]

X = mapping_df.drop(columns=['track_genre_top']).values

# Encode the genre labels
label_encoder = LabelEncoder()
y = label_encoder.fit_transform(mapping_df['track_genre_top']) # Labels matrix

# Check the shapes of the features and labels
print(f"Features shape: {X.shape}")
print(f"Labels shape: {y.shape}")

# Optional: Save the label encoding mapping for later use
genre_mapping = dict(zip(label_encoder.classes_, label_encoder.
# ↪transform(label_encoder.classes_)))
print("Genre mapping:", genre_mapping)
```

Features shape: (7918, 518)

Labels shape: (7918,)

Genre mapping: {'Electronic': 0, 'Experimental': 1, 'Folk': 2, 'Hip-Hop': 3, 'Instrumental': 4, 'International': 5, 'Pop': 6, 'Rock': 7}

```
[7]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import numpy as np

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42, stratify=y)

# Normalize the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Convert data to PyTorch tensors
X_train = torch.tensor(X_train, dtype=torch.float32)
X_test = torch.tensor(X_test, dtype=torch.float32)
y_train = torch.tensor(y_train, dtype=torch.long)
y_test = torch.tensor(y_test, dtype=torch.long)
```

```
[8]: # Define a custom Dataset class
class GenreDataset(Dataset):
    def __init__(self, features, labels):
        self.features = features
        self.labels = labels

    def __len__(self):
        return len(self.labels)

    def __getitem__(self, idx):
        return self.features[idx], self.labels[idx]

# Create Dataset and DataLoader
train_dataset = GenreDataset(X_train, y_train)
test_dataset = GenreDataset(X_test, y_test)

train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)

# Define the neural network
```

```

class GenreClassifier(nn.Module):
    def __init__(self, input_size, num_classes):
        super(GenreClassifier, self).__init__()
        self.fc1 = nn.Linear(input_size, 128)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(128, 64)
        self.fc3 = nn.Linear(64, num_classes)

    def forward(self, x):
        x = self.relu(self.fc1(x))
        x = self.relu(self.fc2(x))
        x = self.fc3(x)
        return x

```

```

[18]: # Initialize the model, loss function, and optimizer
input_size = 518 # Number of input features
num_classes = 8 # Number of classes
model = GenreClassifier(input_size, num_classes=8)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

def train_model(model, train_loader, criterion, optimizer, epochs):
    model.train()
    total_loss = 0
    correct = 0
    total = 0

    for data, labels in train_loader:
        data, labels = data.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(data)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        total_loss += loss.item()
        _, predicted = outputs.max(1)
        correct += (predicted == labels).sum().item()
        total += labels.size(0)

    train_loss = total_loss / len(train_loader)
    train_acc = correct / total
    return train_loss, train_acc

# Testing loop
def test_model(model, test_loader, criterion, device):
    model.eval()

```

```

running_loss = 0.0
correct = 0
total = 0

with torch.no_grad():
    for features, labels in test_loader:
        features, labels = features.to(device), labels.to(device)

        # Forward pass
        outputs = model(features)
        loss = criterion(outputs, labels)

        # Track loss and accuracy
        running_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        correct += (predicted == labels).sum().item()
        total += labels.size(0)

epoch_loss = running_loss / len(test_loader)
epoch_acc = correct / total
return epoch_loss, epoch_acc

```

```

[19]: # Training the model
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model.to(device)

num_epochs = 20
for epoch in range(num_epochs):
    train_loss, train_acc = train_model(model, train_loader, criterion,
    ↪optimizer, epochs=10)
    test_loss, test_acc = test_model(model, test_loader, criterion, device)

    print(f"Epoch {epoch+1}/{num_epochs}")
    print(f"Train Loss: {train_loss:.4f}, Train Accuracy: {train_acc:.4f}")
    print(f"Test Loss: {test_loss:.4f}, Test Accuracy: {test_acc:.4f}")

# Save the trained model
torch.save(model.state_dict(), 'genre_classifier.pth')
print("Model saved to 'genre_classifier.pth'")

```

```

Epoch 1/20
Train Loss: 1.4521, Train Accuracy: 0.4847
Test Loss: 1.2609, Test Accuracy: 0.5612
Epoch 2/20
Train Loss: 1.1181, Train Accuracy: 0.6080
Test Loss: 1.1971, Test Accuracy: 0.5903
Epoch 3/20
Train Loss: 0.9738, Train Accuracy: 0.6514

```

Test Loss: 1.1885, Test Accuracy: 0.5966
Epoch 4/20
Train Loss: 0.8518, Train Accuracy: 0.7024
Test Loss: 1.2233, Test Accuracy: 0.5896
Epoch 5/20
Train Loss: 0.7585, Train Accuracy: 0.7384
Test Loss: 1.2281, Test Accuracy: 0.6061
Epoch 6/20
Train Loss: 0.6513, Train Accuracy: 0.7772
Test Loss: 1.2678, Test Accuracy: 0.5979
Epoch 7/20
Train Loss: 0.5562, Train Accuracy: 0.8118
Test Loss: 1.3424, Test Accuracy: 0.6048
Epoch 8/20
Train Loss: 0.4712, Train Accuracy: 0.8420
Test Loss: 1.4261, Test Accuracy: 0.6067
Epoch 9/20
Train Loss: 0.3850, Train Accuracy: 0.8735
Test Loss: 1.4868, Test Accuracy: 0.5934
Epoch 10/20
Train Loss: 0.3189, Train Accuracy: 0.9020
Test Loss: 1.6245, Test Accuracy: 0.5953
Epoch 11/20
Train Loss: 0.2362, Train Accuracy: 0.9313
Test Loss: 1.7547, Test Accuracy: 0.5922
Epoch 12/20
Train Loss: 0.1807, Train Accuracy: 0.9471
Test Loss: 1.8530, Test Accuracy: 0.5909
Epoch 13/20
Train Loss: 0.1350, Train Accuracy: 0.9651
Test Loss: 2.0017, Test Accuracy: 0.5859
Epoch 14/20
Train Loss: 0.1177, Train Accuracy: 0.9686
Test Loss: 2.1387, Test Accuracy: 0.5859
Epoch 15/20
Train Loss: 0.1009, Train Accuracy: 0.9730
Test Loss: 2.2478, Test Accuracy: 0.5878
Epoch 16/20
Train Loss: 0.0851, Train Accuracy: 0.9779
Test Loss: 2.4011, Test Accuracy: 0.5903
Epoch 17/20
Train Loss: 0.1165, Train Accuracy: 0.9670
Test Loss: 2.6224, Test Accuracy: 0.5789
Epoch 18/20
Train Loss: 0.1336, Train Accuracy: 0.9610
Test Loss: 2.6199, Test Accuracy: 0.5833
Epoch 19/20
Train Loss: 0.0773, Train Accuracy: 0.9769

Test Loss: 2.5974, Test Accuracy: 0.5890
Epoch 20/20
Train Loss: 0.0266, Train Accuracy: 0.9970
Test Loss: 2.6850, Test Accuracy: 0.5909
Model saved to 'genre_classifier.pth'

```
[20]: import random

# Function to test the model with 10 random songs
def test_random_songs(model, X, y, metadata_df, device, label_encoder):
    model.eval() # Set the model to evaluation mode

    # Select 10 random indices from the test dataset
    random_indices = random.sample(range(len(X)), 10)

    # Get the corresponding features, true labels, and metadata
    random_features = X[random_indices]
    true_labels = y[random_indices].numpy()

    # Use the original indices of the test set to retrieve metadata
    metadata = metadata_df.iloc[random_indices]

    # Move features to the appropriate device
    random_features = random_features.to(device)

    # Predict genres
    with torch.no_grad():
        outputs = model(random_features)
        _, predicted_labels = torch.max(outputs, 1)
        predicted_labels = predicted_labels.cpu().numpy() # Move predictions
    ↪back to CPU

    # Decode the true and predicted labels
    true_genres = label_encoder.inverse_transform(true_labels)
    predicted_genres = label_encoder.inverse_transform(predicted_labels)

    # Display the results
    print(f"{'Track ID':<10} {'Title':<30} {'True Genre':<15} {'Predicted',
    ↪Genre':<15}")
    print("-" * 70)
    for i in range(len(random_indices)):
        track_id = metadata.index[i]
        title = metadata.iloc[i]['track_title'] if 'track_title' in metadata.
    ↪columns else "Unknown"
        true_genre = true_genres[i]
        predicted_genre = predicted_genres[i]
```

```
print(f"{track_id:<10} {title:<30} {true_genre:<15} {predicted_genre:
↪<15}")
```

```
# Example usage
```

```
test_random_songs(model, X_test, y_test, tracks_df, device, label_encoder)
```

Track ID	Title	True Genre	Predicted Genre
1154	Hello Heartstring	Pop	Pop
416	Smyrna Snow Walk	Rock	Rock
1082	Nam Nhi-tu	International	International
481	Council Bluffs	Folk	Folk
891	Your One Mind	Rock	Rock
564	The Sugar Society	Electronic	International
1620	Track 01	Experimental	Rock
1525	Scovil	International	International
458	Hunt Like Devil 4	Pop	Folk
1127	Do I Tingle? Up?	Rock	Pop